



UNIVERSITY OF
BIRMINGHAM

SecurelyShare - A Secure Document Sharing System for Android Devices

Julie Sowards

Supervisor: Eike Ritter

Date: September 2014

Title: SecurelyShare: A Secure Document Sharing System for Android Devices

Author: Julie Sowards

Supervisor: Eike Ritter

Date: September 2014

Declaration

The material contained within this dissertation has not previously been submitted for a degree at the University of Birmingham or any other university. The research reported within this dissertation has been conducted by the author unless indicated otherwise.

Signed

Submitted in conformance with the requirements of the post graduate degree of MSc Computer Science at the School of Computer Science, University of Birmingham (2013/2014).

© University of Birmingham 2013/2014

Abstract

This project

Acknowledgements

I would like to express my gratitude to my supervisor Dr. Eike Ritter, whose enthusiasm, support and patience have been so much appreciated throughout this project. Thanks also go to Andrew Walker, Steve Robbins and Josh Wainwright for being there to bounce ideas off, encouraging me to keep going when things got tough and generally keeping me sane and finally to my husband, Phil, without whose unswerving support the last year would not have been possible.

Table of Contents

1	Introduction	2
1.1	Motivation	2
1.2	Project Aims	3
1.3	Overview of Report	3
2	Background Material	5
2.1	Android Security Features	5
2.2	Dropbox API	6
2.3	Review of Related Work	7
2.4	Review of Other Products	8
3	Analysis and Design	10
3.1	Overview	10
3.2	Problem Analysis	11
3.3	Document Storage	12
3.4	Cryptographic Design	13
3.5	Android Prototype	15
4	Implementation	17
4.1	Development Environment	17
4.2	Initial Development	18
4.3	SecurelyShare Prototype	18
4.4	Admin System	21
4.5	Implementation Challenges	21
4.6	Testing	23
5	Project Management	25
6	Evaluation	27
6.1	Security	27
6.2	Product	29
6.3	Process	30
6.4	Overall Project	30
6.5	Future Work	30
7	Conclusion	32

Chapter 1

Introduction

1.1 Motivation

As the popularity of mobile communications devices increases, there is a growing tendency to use these as a convenient means of reviewing and revising documents on-the-move. Where these documents are of a confidential nature, particular attention must be paid to the fact that mobile devices are more vulnerable to compromise than traditional desktops, which are usually more extensively protected by the security measures implemented as part of an organization's internal network.

There are multiple mechanisms for keeping files secure on company servers whilst allowing employees the necessary permissions to work collaboratively with sensitive data as required. As the mobile device culture becomes more prevalent in the workplace, the addition of Mobile Device Management (MDM) applications empowers users to also access corporate data via their mobile devices whilst still allowing IT departments to retain a degree of control over data security.

Thus, it is acknowledged that maintaining the security of confidential documents can be challenging, even with the weight of a corporate IT infrastructure behind it. In this project, we seek to address the issue of allowing groups of users from different organizations (i.e. with no shared IT infrastructure) to collaborate securely on confidential documents and furthermore, to access these documents via a smartphone or tablet computer whilst minimizing the risk of exposing sensitive information to a potential attacker.

1.2 Project Aims

The primary aim of this project was to implement a scheme to facilitate secure sharing of confidential documents between a group of collaborators, subject to the following constraints:-

- Groups are self-organizing and represent multiple organizations, hence they cannot draw on the support of any central IT services.
- The documents involved are confidential in nature and hence should be encrypted both in transit and at rest.
- Group members wish to be able to access documents on a mobile device running the Android operating system (which may involve use of public wi-fi) without compromising security
- The solution devised should use only well-tested cryptographic techniques and standard libraries and should minimize the amount of trust to be placed in a third-party.

Our secondary objective was to gain an understanding of the basic features of the Android platform, explore some of the techniques involved in developing for mobile devices and some of the challenges encountered in working in this type of event-driven environment.

In pursuit of these aims we developed a solution called SecurelyShare, consisting of a detailed design of the security components of the system and a prototype Android application to provide a platform on which to implement and evaluate the various security features. It was acknowledged that, in a live setting, documents would usually originate on a PC rather than on a tablet device and thus the system would also need to a PC-based component. However, within the time constraints of the project it was considered infeasible to develop a fully featured system; our solution is submitted rather as a 'proof of concept'.

1.3 Overview of Report

The subsequent chapters of this report will deal with the design, implementation and evaluation of the project. Chapter 2 introduces some of the background material on key technologies used and presents an overview of the Android applications reviewed as part of our preliminary research. In the light of this research, Chapter 3 presents a detailed analysis of the problem, defines the threat model against which we are attempting to defend, and provides an outline of the final solution design. Chapters 4 and 5 provides details of the implementation, testing and management of the project and finally Chapter 6 evalu-

ates the project components, presents a review of the overall success of project and gives recommendations for future work.

Chapter 2

Background Material

In this chapter we will introduce some key aspects of the android architecture and its security features. We will also examine the features offered by the Dropbox API and finally we will look briefly at some of the commercial applications which were reviewed as part of our initial research and which offer some features similar to Securely Share.

2.1 Android Security Features

Android provides an open source platform and application environment for mobile devices. It has a layer based architecture whose foundational component is the Android Operating System. Based on the tried and tested Linux kernel and modified by Google to include some additional features, it is the comprehensive user permissions model inherent in Linux that is responsible for providing the separation between applications that is a key feature of Android security. At installation time, each application is assigned a unique user ID (UID); at runtime each application is run as a separate process and as a separate user with its given UID. This creates an Application Sandbox, protecting any resources belonging to that application (memory space, files, etc.) from being accessed by another application unless specifically permitted to do so by the developer.

Android also provides some additional security mechanisms. Most relevant to our application is the requirement that each application requiring certain restricted permissions (write to external storage, access network, access contacts, etc.) has to declare them and the user is alerted and asked to confirm this access at installation time. Figure 2.1 provides a useful summary of these mechanisms and the related security issues.

Table 1. Security mechanisms incorporated in Android.

Mechanism	Description	Security issue
Linux mechanisms		
POSIX users	Each application is associated with a different user ID (or UID).	Prevents one application from disturbing another
File access	The application's directory is only available to the application.	Prevents one application from accessing another's files
Environmental features		
Memory management unit (MMU)	Each process is running in its own address space.	Prevents privilege escalation, information disclosure, and denial of service
Type safety	Type safety enforces variable content to adhere to a specific format, both in compiling time and runtime.	Prevents buffer overflows and stack smashing
Mobile carrier security features	Smart phones use SIM cards to authenticate and authorize user identity.	Prevents phone call theft
Android-specific mechanisms		
Application permissions	Each application declares which permission it requires at install time.	Limits application abilities to perform malicious behavior
Component encapsulation	Each component in an application (such as an activity or service) has a visibility level that regulates access to it from other applications (for example, binding to a service).	Prevents one application from disturbing another, or accessing private components or APIs
Signing applications	The developer signs application .apk files, and the package manager verifies them.	Matches and verifies that two applications are from the same source
Dalvik virtual machine	Each application runs in its own virtual machine.	Prevents buffer overflows, remote code execution, and stack smashing

Figure 2.1: Security mechanisms incorporated in Android (Shabtai et al., 2010)

KeyStore

Another of Android's security mechanisms is the KeyStore which is designed for securing cryptographic artefacts. Although frequently associated with the process of digitally signing applications for release, they can more generally be used for storing certificates, private keys, and even symmetric keys. Each KeyStore is secured using Password Based Encryption (PBE) with a user supplied password and may also use separate passwords for each individual entry.

2.2 Dropbox API

Dropbox is a cloud storage service that also offers users automatic backup facilities, file synchronization across devices, and the ability to share files with other users. It provides multi-platform client applications plus a series of public APIs that enable different subsets of the Dropbox functionality to be integrated into third-party applications.

On the Android platform, Dropbox offers three APIs, described on its website as follows:-

Core The Core API includes the most comprehensive functionality including features such as search, file restore, etc. Although more complex than either of the other two to implement, it is often more suitable when developing server-based apps.

Datastore The Datastore API provides a means of storing and synchronizing structured data like contacts, to-do items, and game states across all the user's devices.

Sync The Sync API provides a file system for accessing and writing files to the user's Dropbox. The Sync API manages the process of synchronizing file changes to Dropbox and can also provide the app with notification when changes are made to files stored on the server.

For the purposes of the SecurelyShare app, although it offers the most basic interface to the Dropbox server, the Sync API was deemed to support both the functionality required for the prototype and some additional facilities which could be implemented at a later date in order to improve the overall user experience.

Each app on the Dropbox Platform needs to be registered in the App Console and the developer needs to select which permissions the app requires. These permissions determine the type of data that the app can access in the user's Dropbox. For the Sync API, three levels of permission were available:

1. App folder: this creates a folder in the user's Dropbox with the same name as the app, all files relating to the app are kept here and access is restricted to this folder and its subfolders.
2. File type: the app is given access to the user's entire Dropbox but is restricted only to seeing files of certain types (documents, images, ebooks, etc.)
3. Full Dropbox: the app is given unrestricted access to the user's Dropbox

Developers are also able to submit a request to use a custom file extension with the File type permission.

2.3 Review of Related Work

was intending to include a summary of all the material relating to key distribution that I reviewed - not sure if I am going to be able to do this now

2.4 Review of Other Products

One of the security claims of cloud storage provider, Dropbox, is that user data stored on their servers is fragmented and encrypted using 256-bit AES. However, although users may feel reassured by these claims, it is also to be noted that since this encryption is applied server-side, the servers also have access to the keys required to decrypt this information. Furthermore, the Dropbox privacy policy states that,

"We may disclose your information to third parties if we determine that such disclosure is reasonably necessary to (a) comply with the law; (b) protect any person from death or serious bodily injury; (c) prevent fraud or abuse of Dropbox or our users; or (d) protect Dropbox's property rights."

It may therefore come as little surprise that there are an increasing number of client-side encryption applications available that integrate with Dropbox and other similar cloud services to ensure that, even if files are disclosed, the service providers would have no access to decryption keys. Although there are a significant number of encryption applications available for Android, most were discarded for consideration here as they are primarily designed for a single user to encrypt files prior to cloud storage and had little or no support for key sharing.

When trying to evaluate these products there appeared to be a marked absence of robust reviews from respected members of the security fraternity; many of the review that were available seemed to focus more on the usability of applications rather than on their security features. We chose three commercially produced applications to consider in greater depth, Boxcryptor, Viivo and SafeMonk. These were selected partly because they seemed to attract most mentions (although that may simply be due to the fact that they were commercial products and hence had larger marketing budgets), but also because we were able to find some level of technical description of the security protocols used. There was a common trend among the examined applications to use AES for encryption of user files and some implementation of RSA to manage the key exchange, All three applications used a central server for key management, although each was keen to stress that they had no access to key material that would enable them to decrypt user files.

Boxcryptor and Safemonk both store the user's private key on the server, their justification being that it would allow the user to install the application on multiple devices and they both used password-based encryption (PBE) in order to secure the key on the server. Although the servers use key-stretching algorithms to derive the encryption keys, we should always be aware that users can be very poor at choosing strong passwords and at keeping them

secure. (In a study of password reuse, Bonneau Das et al. (2014) observed that 43 percent of users reused a password across multiple sites)

Chapter 3

Analysis and Design

In this chapter we will present some of the salient points from our initial analysis of the problem, detail the threat model against which we are trying to protect and any significant assumptions made which influenced our design decisions, and finally outline our solution design. It should be noted at this point that we had absolutely no prior experience of developing for (or even using) an Android device, therefore this phase was highly iterative. As our familiarity and understanding of the capabilities of Android platform increased and as the design evolved, we were able to revisit these assumptions and the threat model in order to strengthen the security of the overall system. It is for this reason that we feel that presenting this material in a single chapter is more reflective of the underlying process.

3.1 Overview

At the outset, the following key areas were identified that would need further research:

- finalisation of the threat model and main assumptions
- mechanism to be used for storage and distribution of shared documents
- security design, including any key generation and distribution schemes
- processes involved in setting up or managing groups
- design of the prototype application, including any steps required to protect sensitive key material

3.2 Problem Analysis

Assumptions

As outlined in the introduction, Securely Share is intended for use by small, autonomous groups with a largely static membership. This factor was significant in enabling us to discard a number of complex key distribution protocols which were designed for much larger or more volatile groups. It was also envisaged that groups would have lifetime measured in months, rather than days, hence a small administrative overhead in setting up the group could be considered acceptable in terms of the group duration.

For practical purposes, it is likely that origination of sensitive documents would be carried out on a PC, as much for the practicality of typing as for any security reasons. Therefore we are aware that one or more desktop client applications would be required for any fully-implemented solution. For the purposes of this "proof of concept" project, it was decided that an outline version would be created in Java in order to verify the cryptography, but that time would not permit this to be developed and tested to a level where it could be submitted as part of the finished product.

As discussed in section 3.4 the decision was made to use digital signatures as part of the security design. This generates the requirement for every group member to have a certificate that can be used for authentication. It is left to the discretion of each group to determine whether this should be issued from a recognised Certificate Authority or whether self-signed certificates are to be allowed. (For small groups, the process of manually confirming the certificate fingerprint by telephone may be acceptable, for example.) In order to contain the scope of this project it was necessary to assume that each user either had, or was able to generate a public/private key pair and the associated certificate.

Finally, it was assumed that the user acknowledges the sensitivity of his data and has taken reasonable steps to protect it, including the deployment of a device locking code and strong passwords.

Threat and Trust Model

For the purposes of this project, the following assumption regarding the threat and trust model were made:-

1. a group member may consider all other group members as trusted.

2. the Certificate Authority may be considered as trusted.
3. all data in transit or at rest on a third-party server is considered to be exposed to an attacker. In particular, as well as a direct attack, data on an external server is subject to compromise by an insider within the third-party organisation or, as discussed in 2.4, may be disclosed to comply with legal requests.
4. data at rest on the external storage of the device should always be considered as vulnerable to attack. For example, an attacker may entice the user to install malware on the device(usually from an untrusted source) by presenting it as a new game, etc. - this can then be used to read data from the user's SD card.
5. data at rest within the application's internal storage is protected to some extent by the platform's in-built security mechanisms, although this should still be encrypted to protect against a casual attacker who is able to gain root access for the purposes of snooping in storage. The prototype is not required to implement software protection against a more sophisticated attacker who is able to exploit root access to implement key logging or to steal secrets from within the app memory. However we are aware that the use of hardware-based cryptography via a Trusted Platform Module (TPM) would ameliorate the effect of this type of attack.

3.3 Document Storage

When considering how to manage the storage and sharing of group documents, we initially considered implementing a custom web server. However, it was recognised that this option had a significant IT overhead, both in development and ongoing support and therefore did not fit with our requirements. It was deemed that a better solution would be to leverage the facilities already available via one of the widely used cloud storage services: Dropbox, Google Drive and OneDrive were considered.

Although in a production-level application it would be desirable to allow the user to choose which cloud storage service they wished to use, it quickly became apparent that to implement using multiple APIs was infeasible within the scope of this project. However, as the facilities offered by each were broadly similar, it was felt that developing for a single API in the prototype would be sufficient for the purposes of validating the concept. Dropbox was selected as it is widely used and is well supported in both Windows and Linux environments, the latter being important as this was the development and testing environment being used for the project.

As described in 2.2, Dropbox offers several version of the API for Android and the Sync

API seemed to be the best match for the project requirements. Developers are encouraged to assign each application the least permissions possible, so initially App level permission was chosen. However, it came to light that this did not allow any access to shared folders so the design specification was amended to use the File type permission. This presents a filtered view of the user's Dropbox folder, only permitting the calling application to see files of the specified types. One disadvantage of this is that although all of the user's folders can be seen, various folder operations (e.g. delete, share) are prohibited as the folders may contain hidden files.

When registering SecurelyShare, text and document types were required but it was also necessary to have a file extension to represent our encrypted files. It would have been possible to submit a request to register custom file types with Dropbox but since time was of the essence, we elected to choose two less-commonly used ebook formats as a temporary measure, .xps (for encrypted documents) and .xeb (for encrypted group keys). The .xps extension was simply added to the filename of an encrypted file so that it was easy to recover the original file extension on decryption.

Administration

Although Dropbox provides an excellent environment for collaborative working with its inbuilt mechanisms for file and folder sharing, this does bring with it the initial administrative overhead of creating a group folder and sharing it with all the members. As discussed, the Dropbox Sync API does not currently provide support for this in Android but it was deemed acceptable for our purposes that this process should be completed via the Dropbox website in the usual way.

3.4 Cryptographic Design

There are two main aspects to the cryptographic design as follows:

- the document encryption scheme
- the group key distribution protocol

Document Encryption

As one of the central assumptions is that any third-party server cannot be trusted, it is essential that client-side encryption is used to protect our sensitive data. SecurelyShare uses a symmetric key encryption scheme and is currently implemented using the AES algorithm in Cipher Block Chaining (CBC) mode with a key length of 256 bits. A design requirement was that references to algorithms and key lengths should be included as static variables rather than being hard-coded in order to provide an easy upgrade path or scope to make these user options at a later date.

Since a user of SecurelyShare may be a member of multiple groups at any one time, it is important that the decryption module is able to determine the correct key to use for each document. Initially we thought to capture this information based on the name of the containing folder but quickly realised that this was prone to error as, if a file became disassociated from its correct folder there would be no way of determining where it came from and hence reuniting it with the correct key. It was determined that each file should have a file header written to it made up of an integer containing the length (in bytes) of the group name, a string containing the group name and a fixed length Initialisation Vector - this would provide all the information required to decrypt the file.

Group Key Distribution

By their nature, symmetric schemes require that all parties are in possession of the key, as the same key is used for both encryption and decryption. Where we are just sharing a key between two parties, methods such as the Diffie-Hellman key exchange protocol (Diffie and Hellman, 1976) are widely used. However, the problem of group key distribution is much more challenging and has been the subject of multiple research papers (see 2.3). Although these papers present a number of interesting theoretical schemes with advanced features, our design goal was to produce a practical but secure implementation which matches requirements built with well-established cryptographic primitives and protocols.

Before generating the group key, the group administrator must set up a Java KeyStore containing the private key and corresponding certificate which is to be used to authenticate the group key on distribution. This should be created on the PC to be used for running the Admin module and must be secured with a strong password which is then used by the KeyStore to encrypt the contents using PBE. Each group member should supply a certificate which should then be imported into the KeyStore and finally the key generation module can be run.

A high-level description of the key-generation process is as follows:-

- generate a new symmetric key for the group
- for each user certificate in the KeyStore, encrypt a copy of the group key with the user's public key (as contained in the certificate) using RSA
- sign the encrypted key using the administrator's certificate
- save the encrypted and signed key to the group Dropbox folder

The use of digital signatures is required for each user to be able to validate that the encrypted key they receive is the authentic one and not one substituted by an attacker. This is a vital part of the security protocol and hence has been included in our system design. However, as will be discussed later, due to some of the challenges encountered during the implementation phase, we were not able to implement this aspect of our design fully in the prototype.

Reasons why chose PKI over IBE or password based solution - aim is to use simplest solution that works.

3.5 Android Prototype

Encryption

As outlined in the Problem Analysis (3.2), it is unlikely that large documents will be originated from within SecurelyShare so our design just included a simple edit window for text input which can then be saved to Dropbox as an encrypted file. For completeness, and to aid testing, we also included the ability to encrypt a file which already exists on the device (although this is not to be encouraged as such files may have been exposed to an attacker prior to encryption).

Decryption

Since tablet devices are ideally suited for reading documents, the major requirement of this app is that it is able to take encrypted documents stored in the user's Dropbox, decrypt them with the appropriate symmetric key and present the original text to the user in a readable format. Our requirements necessitated that this should be accomplished without the need to save the plaintext to disk on the device. Initially this seemed as though it would

be achievable as the Android platform provides mechanisms for one application to provide data to another,. However, on deeper study, it transpired that none of these options was exactly right for our purpose. We considered a number of ways of modifying our design in order to overcome this challenge but all of them required some sort of modification to the third-party application. Eventually it was decided that we had to weaken our requirements or risk jeopardising the entire project time scale.

It was agreed that we could achieve our objective for simple text files by displaying the decrypted data in one of the Android-supplied Edittext widgets, which would enable it to be read and, if required, edited saved again in encrypted format. For files requiring a more complex viewing and editing environment, for example those in Portable Document Format (PDF), we would write these temporarily to the external cache before broadcasting an Intent with the appropriate URI in order to allow the user to open this temporary file with one of the PDF readers installed on their device. The security implications of this and suggestions for future work are discussed further in Chapter 6.1.

Key Management

When first installed on the Android device, the SecurelyShare application will create two KeyStores and save them in the application's private storage. The Certificate keystore contains a copy of the user's private key and corresponding certificate chain and the Group keystore is the place where all the group keys will be stored.

As described above, the group key is created by the group administrator and a file is generated for each group member containing a copy of the group key encrypted with the private key contained in the Certificate KeyStore. Before the user can begin creating or reading documents for the group, this key must be imported into the Group keystore on the device. The user selects his file from the Dropbox folder containing the encryption keys for the group, the app decrypts this with the user's private key and stores it in the Group keystore using the group name as an alias.

Chapter 4

Implementation

4.1 Development Environment

An early challenge encountered in this project was setting up a suitable development environment. Initially the intention was to use the Android Development Tools plug-in for Eclipse, as this was an extension to an already familiar IDE, however, difficulties were encountered with the school computers and finally we elected to switch to Android Studio (Beta) as suggested on the Android Developers website (Developer.android.com, 2014a).

Since the main rationale for the Android application was to allow users to read documents away from the PC, we felt that the application should be optimised for use on a tablet. As a difficulty arose in getting the device emulators to work on Linux, it became necessary to purchase a physical device for the purposes of the project. Budget too was a significant consideration; therefore the device chosen was a Hudl 7" tablet running the Android 4.2.2 Jelly Bean operating system. For the purposes of this project, it was decided that we would develop and test for this device. This meant that we were restricted to API 17, which eliminated the need to include any backward compatibility but also precluded using features of the latest Android releases as it transpired that there was no upgrade path available for this device. This fact turned out to be highly significant in a later stage of the project.

4.2 Initial Development

It was felt wise to begin by developing a very simple application in order to gain some familiarity with the complexities of the Android architecture and its API. This initial application enabled us to input some text, encrypt it (using "convert to upper case" in lieu of encryption) and write it to the application storage.

The next stage was to introduce actual encryption, for which we needed to investigate the Java Cryptography Architecture(JCA). Although the choice of algorithm was determined as part of the design, we were mindful of the fact that many security weaknesses are introduced by poor implementation. As Schneier (1998) states,

'And just as it's possible to build a weak structure using strong materials, it's possible to build a weak cryptographic system using strong algorithms and protocols'

For this reason we decided that we would initially develop and test the Java classes to be used for encryption and decryption in the more familiar Linux environment before porting to Android. Linux also provides a simple command line interface which was valuable for inspecting the binary (or hex) content of encrypted files to aid debugging.

4.3 SecurelyShare Prototype

In undertaking this project we were keen to understand more of the tools and techniques of Android development, therefore we attempted to incorporate a number different user interface components including alert and custom dialogs, fragments and ActionBar menus.

Initialisation

The Initialisation Activity is run the first time SecurelyShare is installed. The user is asked whether they wish to import a KeyStore from backup and asked to set up a master password. In the current implementation, the import option must be chosen as the Certificate keystore needs to be created on a PC, either in Linux using `keytool` from the command line or using a program called Portecle (Portecle.sourceforge.net, 2014). An empty keystore is created for the group keys and the user is asked to confirm to proceed. We used




Figure 4.1: Set Master Password

Main Activity

The home page consists simply of four buttons for moving to the other main activities in SecurelyShare. We recognise that this is somewhat visually unappealing and duplicates the menu options included in the Action Bar, however, in line with our design goals, it was felt better to focus on improving security features rather than visual appeal at this stage.

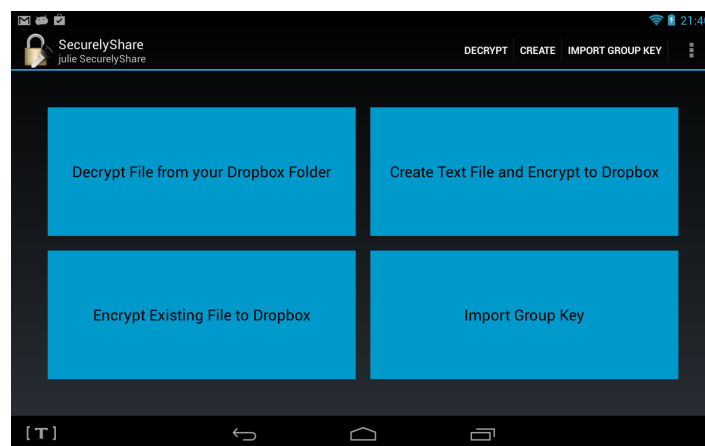


Figure 4.2: Basic home screen

The user is invited to enter the Master password upon opening the application and, having decided that we would not require the password to be re-entered to access each file, it was necessary to ensure that this data was available wherever needed. Although the option of writing the password to the `SharedPreferences` database, from whence it could be accessed by other activities, would have been an easy solution, we did not feel that this represented good practice. Further research revealed this to be a contentious question, with experienced developers divided as to whether this was best done by extending the `Application` class or by use of the Singleton design pattern. As the guidance from Android

Developers website states that 'there is normally no need to subclass Application. In most situation, static singletons can provide the same functionality in a more modular way' (Developer.android.com, 2014b) , we decided to use this approach but recognise we may not have completely understood the full complexities of this problem at this juncture.




Figure 4.3: Dialog to request Master password to unlock access to KeyStores

Other Features

Fragments

Figure 4.4 shows the screen for importing a new group key, which illustrates our use of the Fragment class.

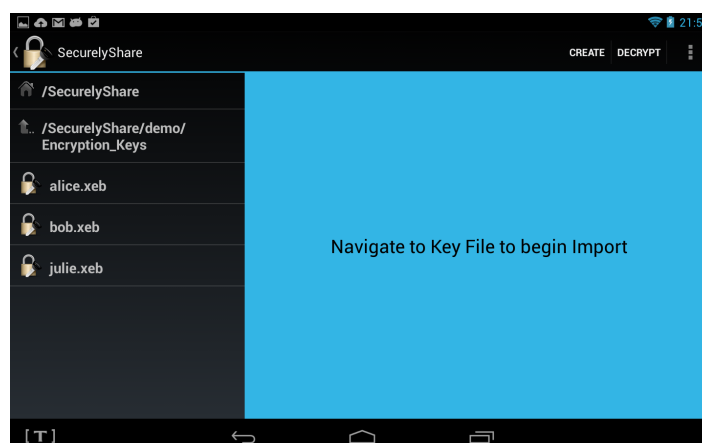


Figure 4.4: Import a group key

Fragments are reusable elements of the application's user interface which can be combined to build a multi-pane user interface (Developer.android.com, 2014c). As shown here, Se-

curelyShare is perhaps best suited to landscape mode as it uses a split-screen, however it could easily be modified to show each fragment in a separate activity on a smaller screen or to improve the display in portrait mode.

Performance and Threading

Although some asynchronous processing results from the calls to the Dropbox API, currently everything in SecurelyShare runs on the main thread. Whilst this may seem an unusual choice when carrying out computationally expensive cryptographic operations, there is little value in introducing parallel processing if it frees the user up to do other tasks. When opening a file, there is little else that the user can practically do apart from wait and in testing, we felt that performance was within acceptable limits. This may need to be revisited in a future iteration; if, for example, a custom PDF processor was implemented where users were needing to encrypt and save large annotated files, it may be preferable to run this on a separate thread.

4.4 Admin System

Admin system – basic with no front end. Designed to run on system with Dropbox installed and running so files uploaded to dropbox simply by saving them in the correct location rather than worrying about using Dropbox API within program. PC version in java that just does basic encryption/decryption. Used primarily for intial development and testing of encryption mechanisms. Single username hardcoded for testing. Fully developed application with user interface outside of scope of this project. Assumption that encrypted blobs are probably also created on PC - simple PC version of program developed to address this, although no gui developed

4.5 Implementation Challenges

Android is highly event-driven; its activities are written using the familiar syntax of Java and, as a new Android developer, it is easy to fall into the trap of thinking of them as 'programs'. However, they could perhaps be more properly thought of as sophisticated event-handlers. Android activities do not control the flow of program execution - they simply provide code to respond to events.

We found that this different way of thinking was one of the most challenging aspects of this project - basic concepts like opening a dialog to ask the user to confirm an action and blocking until the user responds are simply not available. Similarly unfamiliar is the way that the operating system actively manages storage. When the user navigates away from an activity, the operating system holds it in memory in a paused state. If, however, memory resources run short, the current instance of the activity is simply destroyed and a new one created next time the user navigates back to it, causing variables to be reinitialised (although the system saves the state of the activity layout so that it can recreate the screen when the user resumes the application). This uncertainty over the precise state of variables required close attention when coding.

Communication between activities

When wanting to start a new activity, the standard method is to create an `Intent` and then to execute `startActivity(intent)`. Data to be passed to the new activity can be placed in a `Bundle` belonging to the intent. However, in order to pass an object in this way it must be serializable, either using standard Java serialization or by implementing `Parcelable`, Android's own implementation which is lightweight and faster. In `SecurelyShare` we are working with object of the Dropbox API, which are neither `Parcelable` nor serializable. Considerable time was spent investigating whether there was an obvious solution to this which we were simply missing due to inexperience but ultimately it appeared that there was none and we had to design around this limitation.

Security Providers

The Java SE Documentation states:

Reminder: Cryptographic implementations in the JDK are distributed through several different providers ("Sun", "SunJSSE", "SunJCE", "SunRsaSign") for both historical reasons and by the types of services provided. General purpose applications SHOULD NOT request cryptographic services from specific providers. (Oracle, 2014)

Having followed this advice, we discovered that although we were able to decrypt and import group keys in the desktop version of our code, this did not work when ported to Android. This presented a major debugging difficulty as, having isolated the problem, we were essentially comparing two pieces of code that looked identical but performed differently. It transpired that the problem arose because Java and Android use different default providers

and these are not interchangeable. We were able to remedy the problem by installing the Bouncy Castle provider and forcing Java to use that provider.

CHALLENGES Major issues with keystore, certificates No access to key tool in android
THINGS NOT INCLUDED However, as will be discussed in Chapter 4, due to some of the challenges encountered during the implementation phase, we were not able to implement this aspect of our design fully in the prototype. (digital signatures)

4.6 Testing

The encryption and decryption elements of our prototype were largely testing by inspection - known texts were encrypted, the files were checked to ensure that they were not leaking any of the plaintext, and these were in turn decrypted and compared with the original. In general, the nature of encryption means that the main body of a file either decrypts completely or not at all, however particular attention must be paid to the very beginning and end of a file, as incorrect use of the API can lead to errors here. As our project was aimed at document encryption, we tested using a mixture of text (.txt files) and PDFs (although we did also confirm that our processes generalised to other file types). Since randomness plays an important role in ensuring semantic security for symmetric encryption schemes, this made any sort of unit testing extremely difficult. At the outset we considered using fixed initialisation vectors and encryption keys to make our tests repeatable. However, on closer consideration this involved making some significant changes to the program logic relating to the use of the cryptography APIs and so was considered to make this sort of testing meaningless.

Other aspects of the user interface were tested manually, trying to make sure that a range of user inputs and behaviour were covered. However, it should be noted that the prototype was not designed to be releasable code and therefore did not include comprehensive validation of user input. Although certain common scenarios were handled (for example, the user is informed if trying to decrypt a document without first importing the appropriate group key), it was deemed that for the purposes of this prototype it was acceptable that other less common behaviours should be permitted to cause the app to terminate (e.g. if the app permission is revoked via the Dropbox website whilst the app is running).

When a user shifts focus away from the current activity, the activity is kept in memory but is switched into the background. Such processes can be killed at any time by the system in order to reclaim memory. In this event, when the user returns to the activity, it is recreated

by the system - this may cause some expected results if not designed for properly. Since this process is managed by the operating system and is determined by a number of external factors, it can prove difficult to test. If testing in an emulated environment, it is possibly to simulate an artificially small device memory to encourage this behaviour to be triggered. However, as previously mentioned, this option was not available to us so a similar effect was created by rotating the device, which also causes an activity to be recreated. Even though not identical, it proved quite an effective substitute in our testing.

Chapter 5

Project Management

In managing this project we used an iterative approach based on the Agile philosophy, although with such a small development team and no external customer many of the processes and artefacts associated with Agile methodologies were inappropriate. In particular, although we had fully working software at the end of each iteration, it was questionable if it could have been considered a 'potentially shippable product' (Gupta, 2008).

At the outset we had very little knowledge of the Android platform which resulted in it being difficult to form an accurate plan. Although Android applications are generally programmed in Java with which we had some familiarity, we had underestimated how much developing for mobile devices differed from anything we had previously done. We realised that a greater portion of the overall development time needed to be given over to production of the prototype, particularly because one of our initial aims was that we would seek to understand and implement best practice in our code. As a result we had to readjust some of the timescales in our project plan, however this was recognised early and was able to be well managed.

One issue with using an iterative development model was that sometimes time was wasted on learning and development that was critical to get a particular iteration to work but which ultimately ended up being abandoned or didn't contribute significantly to the final product. It was important to guard against retaining features in the final design simply because of the amount of work that had gone into implementing them in an earlier iteration.

During development some elements which were required for security came to light. Some of these were implemented in the prototype, for example we deleted the KeyStore after three failed password attempts, in order to protect against a device which had fallen into the wrong hands, whereas others were merely noted in the security design. Also, as discussed

in Chapter 4 we discovered a difficulty with opening decrypted PDFs without the need to write them out to a temporary file which made it necessary for us to revise the specification of the prototype.

One of the greatest challenges encountered in managing the project was avoiding what is frequently described as "scope creep". It was established at the outset that our development activities should be focussed on maximising the security of the prototype and yet we found it difficult to avoid getting sidetracked into adding features designed to enhance the user interface. The process of delivering a regular progress report to our supervisor proved useful in helping to recognise and correct such digressions.

Chapter 6

Evaluation

6.1 Security

In designing our application we applied the principle of 'Defense in Depth' (NSA, 2000), using multiple layers of protection to guard against failure in any single mechanism.

- linux and android inbuilt mechanism provide sandboxing and control access to internal file storage
- the device should be locked with a passcode
- the application has a master password that is used to encrypt the KeyStore containing all the cryptographic artefacts
- data stored on both the device and Dropbox is encrypted
- encryption is performed on the device to ensure that the server has no access to decryption keys
- using the Dropbox Sync API automatically ensures that all data in transit is protected using SSL
- data on the server is also protected by the Dropbox access control facilities

Dropbox uses SSL so protected against replay attacks

Client side encryption was used because of the desire to ensure that server had no access to plaintext

Refer to examples which use PBE. Happy to use it for protection of key material on device as even with password, attacker would have to have access to device or would require more sophisticated remote attack which is out of our scope.

Unauthorised Access to device

loss of device - we need to include some protection against loss of device but an attacker who is able to gain access to the device whilst the app is running would be able to use the app to decrypt any documents in Dropbox as key store would be unlocked. It should be noted however that the use of Dropbox does provide some additional protection for this as, once aware that the device is lost, the user is able to log in via the Dropbox website and revoke SecurelyShare's permissions for the account with immediate effect.

Password Policy

Password policy - balance of requiring password for every file access which encourages user to choose insecure password. Opted to use one password both to unlock the KeyStore and to encrypt each of the group keys. Could have had a separate password for each group but requiring user to remember different passwords often leads to insecure ones.

Keystore deleted

Key Distribution

Each group requires an administrator to carry out the initial setup, although any member can serve in this role. It is also possible to delegate this to someone who is not part of the group without giving them access to the group encryption key. There is a recognised weakness in this solution if a corrupt administrator imports the attacker's certificate into the KeyStore, thus generating a copy of the encryption key for the attacker. After this point there would be no benefit for an attacker in subverting them, as being able to get access to the KeyStore containing all the certificates wouldn't give any advantage. Gaining access to the Dropbox folder would enable the attacker to see what files were in there but without private keys he wouldn't be able to decrypt anything.

Backward Secrecy

This is often required in a messaging environment where it is important that new group members are not able to decrypt messages from before they joined the group. However, in our application backward secrecy is not required as all group members should have access to historic documents.

Forward Secrecy and Key Revocation

Requiring perfect forward secrecy adds an additional level of complexity to any encryption. The main application of forward secrecy to our project is that a member who leaves the group should not have access to future documents. Although our current cryptographic design does not provide this (once established, the group key is never changed), one of the benefits of using a service like Dropbox is that we are able to avail ourselves of its extensive access control mechanisms. In order to compromise security an attacker needs access to both the group's shared Dropbox folder *and* the group key - a member leaving the group would only have access to the former, which is considered sufficient for our purposes at this point. Similarly, there is currently no means of revoking a key without the need to re-key all of the documents for the group (although again we could revoke access to the folder).

This is one of the reasons why we suggest our solution may not be scalable to large groups as in this instance the membership is much more likely to be volatile. An evaluation of more sophisticated key exchange protocols in order to strengthen this aspect of security without introducing dependency on a server for key management could be the subject of further work.

Known Vulnerabilities

Had to compromise on aim never to write to storage - this is a temporary measure but would need to implement a custom pdf reader which is beyond the scope of this project.

digital signatures to protect against man in the middle

6.2 Product

scalability - Structure would work with larger group but there may be better key management protocols - admin issue with setting up dropbox shares etc.

PROTOTYPE EVALUATION dependent on exactly correct alias for groupid and folder name • Is designed as a "proof of concept" • Aspires to use "best practice" within the code • Uses well-tested cryptographic techniques and standard libraries • Adheres to the stated security requirements No ability to change passwords etc added at present

6.3 Process

major challenge of the fact that android is an operating system not a programming language
- event driven programming

Since many of the example apps that we looked at were focused on personal encryption with the sharing capacity as an added feature, it was important that we remembered that the core purpose of our application was group collaboration and therefore central to any design we might choose was the requirement for a robust key sharing mechanism.

development environment

6.4 Overall Project

EXPLAIN HOW WELL SOLUTION MEETS OBJECTIVES ACHIEVEMENTS

WHAT YOU HAVE LEARNED

- Android is a whole new operating system not just 'Java with extra bits'

Problem with API meant that we had to generate keystore for private key and corresponding certificate chain manually. Further work would be to redesign this part of the app using the additional features offered in Android 4.3. Currently consider this to be one of the weakest parts of the design - by the time I realised that this could have been done better, it was too late to go back and change it.

6.5 Future Work

look at generating keystores on device - would need to include method of exporting keystore (encrypted?) for distribution to other device.

Review other schemes - lots of theoretical examples but anything that would work in practice

Implement proper user interface

Write a proper PC-based client application and look at improving group names

different levels of password policy as a user option - this may be a useful addition to be considered in the future. Similarly, for ease of testing and general usability we decided not to implement encryption of filenames in the prototype, However, this would again be a potential user option.

nelenkov.blogspot.co.uk - credential storage enhancements in Android 4.3

Chapter 7

Conclusion

was not in innovative design and algorithms but in the correct deployment of the Dropbox and cryptographic libraries, and in managing the flow of information between the various activities

Bibliography

- boxcryptor. Technical overview — how boxcryptor works — aes and rsa. <https://www.boxcryptor.com/en/technical-overview>, 2014. Accessed: 2014-10-05.
- K. Cairns and G. Steel. Developer-resistant cryptography. <https://www.w3.org/2014/strint/papers/48.pdf>, 2014. Accessed: 2014-10-05.
- Y. Challal and H. Seba. Group key management protocols: A novel taxonomy. *International Journal of Information Technology*, 2(1):105–118, 2005.
- R. Cojocaru. Protecting stored data on android devices. *Journal of Mobile, Embedded and Distributed Systems*, 6(1), 2014. URL http://jmeds.eu/index.php/jmeds/article/view/Protecting_Stored_Data_on_Android_Devices.
- A. Das, J. Bonneau, M. Caesar, N. Borisov, and X. Wang. The tangled web of password reuse. In *Proceedings of NDSS*, 2014.
- Developer.android.com. Android studio — android developers. <http://developer.android.com/sdk/installing/studio.html>, 2014a. Accessed: 2014-10-03.
- Developer.android.com. Application — android developers. <http://developer.android.com/reference/android/app/Application.html>, 2014b. Accessed: 2014-10-03.
- Developer.android.com. Fragments — android developers. <http://developer.android.com/guide/components/fragments.html>, 2014c. Accessed: 2014-10-03.
- W. Diffie and M. E. Hellman. New directions in cryptography. *Information Theory, IEEE Transactions on*, 22(6):644–654, 1976.
- Dropbox. Dropbox platform developer guide. <https://www.dropbox.com/developers/reference/devguide>, 2014. Accessed: 2014-10-05.
- M. Fowler and J. Highsmith. The agile manifesto. *Software Development*, 9(8):28–35, 2001.

- A. Franz and T. Brants. All our n-gram are belong to you. <http://googleresearch.blogspot.co.uk/2006/08/all-our-n-gram-are-belong-to-you.html>, 2006. Accessed: 2014-10-05.
- Google. Android security program overview. <http://source.android.com/devices/tech/security/#android-platform-security-architecture>, 2014. Accessed: 2014-10-05.
- D. Grolimund, L. Meisser, S. Schmid, and R. Wattenhofer. Cryptree: A folder tree structure for cryptographic file systems. In *Reliable Distributed Systems, 2006. SRDS'06. 25th IEEE Symposium on*, pages 189–198. IEEE, 2006.
- S. Gunasekera. *Android Apps Security*. Apress, Berkely, CA, USA, 1st edition, 2012. ISBN 1430240628, 9781430240624.
- M. Gupta. Definition of done: A reference. <https://www.scrumalliance.org/community/articles/2008/september/definition-of-done-a-reference>, 2008. Accessed: 2014-10-05.
- H. Harney. Group key management protocol (gkmp) architecture. *Group*, 1997.
- NSA. Defense in depth. https://www.nsa.gov/ia/_files/support/defenseindepth.pdf, 2000. Accessed: 2014-10-03.
- Oracle. Java cryptography architecture oracle providers documentation for java platform standard edition 7. <http://docs.oracle.com/javase/7/docs/technotes/guides/security/SunProviders.html>, 2014. Accessed: 2014-10-05.
- A. Perrig. Efficient collaborative key management protocols for secure autonomous group communication. In *International Workshop on Cryptographic Techniques and E-Commerce (CrypTEC'99)*, pages 192–202, 1999.
- Portecle.sourceforge.net. Portecle: Home, 2014. URL <http://portecle.sourceforge.net/>.
- R. V. Rao, K. Selvamani, and R. Elakkiya. A secure key transfer protocol for group communication. *arXiv preprint arXiv:1212.2720*, 2012.
- SafeMonk. Safemonk is serious about encryption key security. <https://www.safemonk.com/security>, 2014. Accessed: 2014-10-05.
- B. Schneier. Security pitfalls in cryptography. In *EDI FORUM-OAK PARK-*, volume 11, pages 65–69. THE EDI GROUP, LTD., 1998.

- A. Shabtai, Y. Fledel, U. Kanonov, Y. Elovici, S. Dolev, and C. Glezer. Google android: A comprehensive security assessment. *IEEE Security and Privacy*, 8(2):35–44, 2010.
- J. Six. *Application Security for the Android Platform: Processes, Permissions, and Other Safeguards*. " O'Reilly Media, Inc.", 2011.
- Y. Zhou, X. Zhang, X. Jiang, and V. W. Freeh. Taming information-stealing smartphone applications (on android). In *Trust and Trustworthy Computing*, pages 93–107. Springer, 2011.